



## Outline

Pandas is a powerful and versatile open-source data analysis and manipulation library for Python. It provides data structures and functions that make it easy to work with structured data, enabling users to perform complex data transformations and analyses efficiently.

Pandas is widely used in data science, machine learning, and statistical analysis due to its ease of use, flexibility, and the ability to handle large datasets efficiently. Whether-

we're analyzing financial data, conducting scientific research, or working on machine learning projects, Pandas is an essential tool in the Python ecosystem.

In this lecture, we will explore the different features of Pandas library, including data structures and data manipulation.

## Getting started

Using the Pandas library requires the version of Python to be 3.5 and above. We can install this library using the following command:

```
pip3 install pandas
```

After installing the library, we must import Pandas, in order to use it in our program.

```
import pandas as pd
```

## Series and Dataframes

A **Series** in Pandas is similar to any other series we come across. A series is a one-dimensional labeled array that can hold any data type.

```
series = pd.Series([2, -8, 3, 9], index=['w', 'x', 'y', 'z'])
```

w	2
x	-8
y	3
z	9

A **Dataframe** in Pandas library is a two-dimensional labeled data structure where each row represents an observation.

```
df = pd.DataFrame(  
    {"words" : ['good', 'better', 'best'],  
     "number" : [22, 33, 44],  
     "names" : ['one', 'two', 'three']},      index = [1, 2, 3])
```

	words	number	names
1	good	22	one
2	better	33	two
3	best	44	three

## Working with Dataframes

In order to read a CSV file in Pandas, we use the ***pd.read\_csv()*** function.

```
pd.read_csv('files.csv')
```

We can take a look at the first two lines of a dataframe.

```
df.head(2)
```

	words	number	names
1	good	22	one
2	better	33	two

We can also take a look at the rows between certain indices.

```
df[3:5]
```

## Working with Dataframes

We can take a look at the last two lines in a dataframe using the ***df.tail()*** function.

```
df.tail(2)
```

	words	number	names
2	better	33	two
3	best	44	three

We can see the analysis of numerical columns in a dataframe using the ***df.describe()*** function.

```
df.describe()
```

	number
count	3.0
mean	33.0
std	11.0
min	22.0
25%	27.5
50%	33.0
75%	38.5
max	44.0

## Working with Dataframes

We can take a look at only the columns in a dataframe.

```
df.columns
```

```
Index(['#', 'Name', 'Type 1', 'Type 2', 'HP', 'Attack', 'Defense', 'Sp. Atk',  
      'Sp. Def', 'Speed', 'Generation', 'Legendary'],  
      dtype='object')
```

We can also see specific columns in a dataframe. Note that instead of the *Name* keyword, we can use the other names of the columns of the CSV file.

```
df['Name']
```

```
0          Bulbasaur  
1          Ivysaur  
2          Venusaur  
3  VenusaurMega Venusaur  
4          Charmander  
...  
795          Diancie  
796  DiancieMega Diancie  
797  HoopaHoopa Confined  
798  HoopaHoopa Unbound  
799          Volcanion
```

## Working with Dataframes

We can take a look at more than two columns at the same time.

```
df[['Name', 'HP']]
```

	Name	HP
0	Bulbasaur	45
1	Ivysaur	60
2	Venusaur	80
3	VenusaurMega Venusaur	80
4	Charmander	39
...	...	...
795	Diancie	50
796	DiancieMega Diancie	50
797	Hoopahoop Confined	80
798	Hoopahoop Unbound	80
799	Volcanion	80

We can also see a specific number of names. Note that we can put any other number in place of 5.

```
df['Name'][5:0]
```

```
0      Bulbasaur
1      Ivysaur
2      Venusaur
3  VenusaurMega Venusaur
4      Charmander
```



## Working with Dataframes

We can take a look at one specific row using the *iloc* keyword.

```
df.iloc[2]
```

```
#          3
Name      Venusaur
Type 1     Grass
Type 2     Poison
HP         80
Attack     82
Defense    83
Sp. Atk    100
Sp. Def    100
Speed      80
Generation 1
Legendary  False
```

We can also see multiple rows of data at the same time.

```
df.iloc[1:4]
```

	#	Name	Type 1	Type 2	HP	Attack	Defense	Sp. Atk	Sp. Def	Speed	Generation	Legendary
3	3	VenusaurMega Venusaur	Grass	Poison	80	100	123	122	120	80	1	False
4	4	Charmander	Fire	NaN	39	52	43	60	50	65	1	False

It is also possible to see the value of a specific row and column.

```
df.iloc[5,1]
```

```
'Charmeleon'
```

## Working with Dataframes

We can observe all the info of the dataframe using the *info()* function.

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 800 entries, 0 to 799
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   #           800 non-null   int64
1   Name        800 non-null   object
2   Type 1      800 non-null   object
3   Type 2      414 non-null   object
4   HP          800 non-null   int64
5   Attack      800 non-null   int64
6   Defense     800 non-null   int64
7   Sp. Atk     800 non-null   int64
8   Sp. Def     800 non-null   int64
9   Speed       800 non-null   int64
10  Generation  800 non-null   int64
11  Legendary   800 non-null   bool
12  total       800 non-null   int64
13  Total       800 non-null   int64
dtypes: bool(1), int64(10), object(3)
memory usage: 82.2+ KB
```

## Exporting Dataframes

We can export a dataframe into an excel file using the ***to\_excel()*** function. We can also export a dataframe into a text file using the ***to\_csv()*** function. It is also possible to export a dataframe into a CSV file, using the ***to\_csv()*** function. It is important to note that when the script is executed, the new file will be located within the project's root directory. If we do not want to include the index numbers, we can set the *index* keyword to False.

```
df.to_excel('file.xlsx')
```

```
df.to_csv('file.txt')
```

```
df.to_csv('file.csv')
```

```
df.to_csv('files.csv' , index=False)
```

## Sorting Dataframes

We can sort the data by alphabetical name in ascending or descending order, using the ***df.sort\_values()*** function. The default order is ascending. However, If we want to sort in descending order, we set the *ascending* keyword to False.

```
import pandas as pd

# Create a sample DataFrame
df = pd.DataFrame({
    'student': ['monica', 'nathalia', 'anastasia', 'marina', 'ema'],
    'grade': ['excellent', 'excellent', 'good', 'very good', 'good']
})

# Sort the DataFrame by the 'student' column in ascending order
df_sorted = df.sort_values(by='student')

# Display the sorted DataFrame
print(df_sorted)
```

	student	grade
2	anastasia	good
4	ema	good
3	marina	very good
0	monica	excellent
1	nathalia	excellent

```
import pandas as pd

# Create a sample DataFrame
df = pd.DataFrame({
    'student': ['monica', 'nathalia', 'anastasia', 'marina', 'ema'],
    'grade': ['excellent', 'excellent', 'good', 'very good', 'good']
})

# Sort the DataFrame by the 'student' column in descending order
df_sorted = df.sort_values(by='student', ascending=False)

# Display the sorted DataFrame
print(df_sorted)
```

## Applying changes to a dataframe

We can add a new column to the dataframe.

```
import pandas as pd

# Create a sample DataFrame
df = pd.DataFrame({
    'student': ['monica', 'nathalia', 'anastasia', 'marina', 'ema'],
    'grade': ['excellent', 'excellent', 'good', 'very good', 'good']
})

# Add a new column 'age'
df['age'] = [21, 22, 23, 24, 25]

print(df)
```

	student	grade	age
0	monica	excellent	21
1	nathalia	excellent	22
2	anastasia	good	23
3	marina	very good	24
4	ema	good	25

We can also add a new column to the dataframe which is a combination of other columns.

```
import pandas as pd

# Create a sample DataFrame
df = pd.DataFrame({
    'student': ['monica', 'nathalia', 'anastasia', 'marina', 'ema'],
    'math_score': [85, 90, 78, 88, 92],
    'science_score': [80, 85, 88, 90, 95]
})

# Add a new column 'total_score' which is the sum of 'math_score' and 'science_score'
df['total_score'] = df['math_score'] + df['science_score']

# Display the updated DataFrame
print(df)
```

	student	math_score	science_score	total_score
0	monica	85	80	165
1	nathalia	90	85	175
2	anastasia	78	88	166
3	marina	88	90	178
4	ema	92	95	187

## Applying changes to a dataframe

We can move the place of a column in a dataframe.

```
import pandas as pd

# Create a sample DataFrame
df = pd.DataFrame({
    'student': ['monica', 'nathalia', 'anastasia', 'marina', 'ema'],
    'math_score': [85, 90, 78, 88, 92],
    'science_score': [80, 85, 88, 90, 95],
    'history_score': [75, 80, 85, 90, 95]
})

# Display the original DataFrame
print("Original DataFrame:")
print(df)

# Get the current columns
cols = df.columns.tolist()

# Move 'math_score' to the end
df = df[cols[0:1] + cols[2:] + [cols[1]]]

# Display the updated DataFrame
print("\nUpdated DataFrame with 'math_score' moved to the end:")
print(df)
```

Original DataFrame:

	student	math_score	science_score	history_score
0	monica	85	80	75
1	nathalia	90	85	80
2	anastasia	78	88	85
3	marina	88	90	90
4	ema	92	95	95

Updated DataFrame with 'math\_score' moved to the end:

	student	science_score	history_score	math_score
0	monica	80	75	85
1	nathalia	85	80	90
2	anastasia	88	85	78
3	marina	90	90	88
4	ema	95	95	92

We can also change the value at a specified row and column.

```
# Change the value of 'math_score' for 'nathalia' (row index 1)
df.loc[1, 'math_score'] = 95 # Using loc
```

## Applying changes to a dataframe

We can apply conditional change to a dataframe.

```
import pandas as pd

# Create a sample DataFrame
df = pd.DataFrame({
    'student': ['monica', 'nathalia', 'anastasia', 'marina', 'ema'],
    'math_score': [85, 90, 78, 88, 92],
    'science_score': [80, 85, 88, 90, 95]
})

# Display the original DataFrame
print("Original DataFrame:")
print(df)

# Increase math_score by 5 for students with a score below 80
df.loc[df['math_score'] < 80, 'math_score'] += 5

# Display the updated DataFrame
print("\nUpdated DataFrame after conditional change:")
print(df)
```

Original DataFrame:

	student	math_score	science_score
0	monica	85	80
1	nathalia	90	85
2	anastasia	78	88
3	marina	88	90
4	ema	92	95

Updated DataFrame after conditional change:

	student	math_score	science_score
0	monica	85	80
1	nathalia	90	85
2	anastasia	83	88
3	marina	88	90
4	ema	92	95

# Changed from 78 to 83

## Filtering data in a dataframe

We can get the rows with a matching value at a specified column.

```
# Get rows where math_score is equal to 90 using loc
matching_value = 90
matching_rows = df.loc[df['math_score'] == matching_value]

# Display the matching rows
print("\nRows with matching math_score value of 90:")
print(matching_rows)
```

```
Rows with matching math_score value of 90:
   student  math_score  science_score
1  nathalia         90             85
3   marina         90             90
```

We can get the rows with multiple matching value at a specified column.

```
# Get rows where math_score is either 90 or 92 using loc and isin
matching_values = [90, 92]
matching_rows = df.loc[df['math_score'].isin(matching_values)]

# Display the matching rows
print("\nRows with matching math_score value of 90 or 92:")
print(matching_rows)
```

```
Rows with matching math_score value of 90 or 92:
   student  math_score  science_score
1  nathalia         90             85
3   marina         90             90
4     ema         92             95
```



## Filtering data in a dataframe

We can retrieve rows from a dataframe where the values in a specified column exceed a certain threshold.

```
# Specify the threshold value
threshold_value = 85

# Get rows where math_score is greater than the threshold value using loc
matching_rows = df.loc[df['math_score'] > threshold_value]

# Display the matching rows
print("\nRows with math_score greater than 85:")
print(matching_rows)
```

	student	math_score	science_score
1	nathalia	90	85
3	marina	88	90
4	ema	92	95

We can also filter rows in a dataframe based on whether a specific column contains a certain substring or word.

```
# Specify the word to filter by
word_to_filter = 'na'

# Get rows where the 'student' column contains the specified word using loc
and str.contains
matching_rows = df.loc[df['student'].str.contains(word_to_filter,
case=False)]

# Display the matching rows
print("\nRows where 'student' contains the word 'na':")
print(matching_rows)
```

	student	math_score	science_score
1	nathalia	90	85
2	anastasia	78	88
3	marina	88	90

## Counting records in a dataframe

```
import pandas as pd

# Create a sample DataFrame
data = {
    'student': ['monica', 'nathalia', 'anastasia', 'marina', 'ema',
'monica', 'nathalia'],
    'grade': ['A', 'B', 'A', 'C', 'B', 'A', 'B']
}

df = pd.DataFrame(data)

# Display the original DataFrame
print("Original DataFrame:")
print(df)

# Use groupby and count to count the number of records for each distinct
value in the 'grade' column
grade_counts = df.groupby('grade').count()

# Display the counts of each distinct value
print("\nCount of each distinct value in the 'grade' column using groupby
and count:")
print(grade_counts)
```

We can count the number of records for each distinct value in a specific column of a Pandas dataframe.

Original DataFrame:

	student	grade
0	monica	A
1	nathalia	B
2	anastasia	A
3	marina	C
4	ema	B
5	monica	A
6	nathalia	B

Count of each distinct value in the 'grade' column using groupby and count:

	student	grade
grade		
A	3	3
B	4	4
C	1	1